

Project 1

Zack Owens

6/9/2023

## 1.1

```
In [4]: eps32 = eps(0f0)
```

```
1.0f-45
```

```
In [5]: eps64 = eps(0.0)
```

```
5.0e-324
```

Meanings of these values: The method `eps()` according to the julia documentation is the distance between 1.0 and the next largest floating-point value of `Float64`. [Source \(https://docs.julialang.org/en/v1/manual/integers-and-floating-point-numbers/\)](https://docs.julialang.org/en/v1/manual/integers-and-floating-point-numbers/)

Chat GPT answer:

In Julia, `eps(x)` returns the difference between `x` and the next representable value in the floating-point number system. The behavior you described, where `eps(0f0)` and `eps(0)` yield different values, can be explained by the fact that `0f0` and `0` are represented as different types of floating-point numbers.

In Julia, the `f0` suffix is used to denote single-precision floating-point numbers, also known as `Float32`, while the default floating-point type is double-precision, known as `Float64`.

The reason `eps(0f0)` and `eps(0)` yield different results is because they are based on the precision of their respective floating-point types. The `Float32` type has lower precision compared to `Float64`, meaning it can represent smaller differences between values. Consequently, the value returned by `eps(0f0)` will be smaller than the value returned by `eps(0)`.

```
In [11]: typeof(0f0)
```

```
(Float32, Inf32)
```

```
In [7]: typeof(0.0)
```

```
Float64
```

## 1.2

```
In [ ]: function sum_values_does_not_work(v)
        sum = 0.0
        for i in v
            sum += i
        end
        return sum
    end
```

```
sum_values (generic function with 1 method)
```

```
In [23]: @time sum_values(v)
```

```
2.419582 seconds (7.53 k allocations: 510.795 KiB, 0.65% compilation time)
2.802596928649634e-36
```

```
In [28]: @time foldr(+,v)
```

```
2.612031 seconds (4 allocations: 64 bytes)
2.3509887f-38
```

```
In [27]: @time foldr(+,v,init=0.0)
```

```
2.552371 seconds (33.84 k allocations: 2.340 MiB, 1.89% compilation time)
2.802596928649634e-36
```

```
In [20]: @time reduce(+,v)
```

```
0.633802 seconds (108.92 k allocations: 7.286 MiB, 9.50% compilation time)
2.802597f-36
```

```
In [21]: @time sum(v)
```

```
0.584358 seconds (33.30 k allocations: 2.195 MiB, 5.93% compilation time)
2.802597f-36
```

The values above can be summarized by looking at the speeds and number of allocations for each function. The sum function performed the best with a correct solution. This function only took 0.584358 seconds with only 2.195 Megabytes of memory allocation. The next fastest function was the reduce function which took 0.633802 seconds with only 7.286 Megabytes of allocation. This method also created the correct answer. The foldr method was more than 4 times slower than the reduce and sum functions and performed the worst. This method took 2.552371 with 2.340 Megabytes of allocations. This method also did not create the correct answer unless the init paramters was specified. The for loop function performed the second worst out of all the fuctions and only returned the correct answer when sum was intialized at value 0.0. This function completed in 2.419582 seconds with less allocations than reduce and sum.

The for loop and foldr did not create the correct answers unless the inital values specifed was a float32 instead julia assuming the answer is a float16 (a 0 or 0f0).

## 2.1

```
In [25]: fib_1(n) = first([1 1;1 0]^(n-1)*[1,0])
         fib_2(n) = last( [1 1;1 0]^(n-1)*[1,1])
```

```
fib_2 (generic function with 1 method)
```

```
In [34]: [1 1; 1 0]^(3)
```

```
2×2 Matrix{Int64}:  
 3  2  
 2  1
```

The array holds the state for the fibonacci sequence. By exponentiating the array you are creating the sequence with the previous values populating different elements in the array. The final multiplication will yield the final number in the first element of the matrix or the last depending on the final multiplying array.

Chat GPT:

(n-1) subtracts 1 from the input number n. This is because the matrix representation assumes a 0-based index for Fibonacci numbers ( $F(0) = 0$ ,  $F(1) = 1$ ,  $F(2) = 1$ , etc.), while n may be given as a 1-based index.

$^$  is the exponentiation operator in Julia. It raises the matrix  $[1 \ 1; 1 \ 0]$  to the power of (n-1). This operation involves matrix exponentiation, as explained earlier.

- $*$  is the matrix multiplication operator in Julia. It multiplies the result of the exponentiation by the column vector  $[1, 0]$ .

Finally, `first()` returns the first element of the resulting vector, which corresponds to the Fibonacci number  $F(n)$ .

## 2.2

```
In [26]: function fib_rec(n)  
          if(n==0 || n==1)  
              return 1  
          else  
              return fib_rec(n-1) + fib_rec(n-2)  
          end  
      end
```

```
fib_rec (generic function with 1 method)
```

## 2.3

```
In [32]: @time fib_1(44)
```

```
0.000009 seconds (11 allocations: 1024 bytes)
```

```
701408733
```

```
In [31]: @time fib_2(44)
```

```
0.000005 seconds (11 allocations: 1024 bytes)
```

```
701408733
```

```
In [30]: @time fib_rec(44)
```

```
4.172246 seconds
```

```
1134903170
```

Overall the recursive function is the worst performing as it requires multiples calculations of the same function call as there is no dynamic programming simplification. While the fib\_1 and fib\_2 work in linear time as the variables are store in matrix. This causes there to not be recursive calls to the same function. Instead this function store the last values in the matrix. The fib\_2 is the fastest taking 0.000005 seconds to run. fib\_1 has a similar runtime as fib\_2 but slightly slower at 0.000009. The fib\_rec function is the slowest function by far as it takes 4.172246 seconds this is very slow compared to the other fib functions.

## 3.1

```
In [5]: function sinx(x, n = 5; degree = true)
```

```
    if degree
```

```
        xRad = deg2rad(x % 360)
```

```
    else
```

```
        xRad = x
```

```
    end
```

```
    value = xRad
```

```
    numer = xRad^3.0
```

```
    denom = 6.0
```

```
    sign = -1.0
```

```
    for i in 1:n
```

```
        value += sign*numer/denom
```

```
        numer *= xRad^2.0
```

```
        denom *= (2.0 * i + 2.0) * (2.0 * i + 3.0)
```

```
        sign *= -1.0
```

```
    end
```

```
    return value
```

```
end
```

```
sinx (generic function with 2 methods)
```

```
In [153]: sum = 0.0
for i in 0:1:90
    sum += abs(sinx(i,6) - sind(i))
end
valueToBeat = 0.1e-7
if sum < valueToBeat
    print("N is within acceptable range")
end
```

N is within acceptable range

The sinx function will convert x to radians if the degrees paramter is true. The value n represents the number of terms that is used in the taylor seires to generate the output value. The larger the n value the smaller the difference between actual sin value and the calculated value.

The second code block sums the absolute differences between the sinx function and the sind function each of these differences are then added to the sum. If the sum is less than the value to beat which was specifically 0.1e7 then the program block will print out "N is within acceptable range". The smallest accepted value is n=6.

## 3.2

```
In [162]: values = [sinx(x,6) for x in 0:10:90]
pdegrees = rpad("degrees",12)
pvalue = rpad("value",12)
printstyled(pdegrees," ", pvalue, " ", color=:yellow)
println()
j = 1
for i in 0:10:90
    pdegree = rpad(i,12)
    pvalue = rpad(values[j],12)
    println(pdegree,pvalue)
    j +=1
end
```

| degrees | value               |
|---------|---------------------|
| 0       | 0.0                 |
| 10      | 0.17364817766693033 |
| 20      | 0.34202014332566877 |
| 30      | 0.5                 |
| 40      | 0.6427876096865427  |
| 50      | 0.7660444431190767  |
| 60      | 0.8660254037859597  |
| 70      | 0.9396926208012458  |
| 80      | 0.984807753125684   |
| 90      | 1.0000000006627803  |

```
In [17]: function runSinx(inc)
        for k in inc
            sinx(k,10)
        end
    end

    function runSind(inc)
        for k in inc
            sind(k)
        end
    end

    function timeSins(inc)
        pinc = rpad("increment",12)
        psinx = rpad("sinx time",12)
        psind = rpad("sind time",12)
        printstyled(pinc," ", psinx, " ", psind, "\n",color=:yellow)

        for i in inc
            sinxTime = @elapsed runSinx(1.0:i:90.0)
            sindTime = @elapsed runSind(1.0:i:90.0)
            println(rpad("0:$i:90",12),rpad(sinxTime,12),rpad(sindTime,13))
        end
    end
end
```

timeSins (generic function with 1 method)

```
In [18]: inc = [0.1, 0.01, 0.001, 0.000_1, 0.000_01, 0.000_001]
        timeSins(inc)
```

| increment   | sinx time | sind time |
|-------------|-----------|-----------|
| 0:0.1:90    | 6.35e-5   | 1.79e-5   |
| 0:0.01:90   | 0.0005987 | 0.0001792 |
| 0:0.001:90  | 0.0066951 | 0.0015582 |
| 0:0.0001:90 | 0.0593895 | 0.0164536 |
| 0:1.0e-5:90 | 0.5756526 | 0.169668  |
| 0:1.0e-6:90 | 5.7955757 | 1.8052115 |

Overall sind function is always faster than the sinx function. Usually the sinx function takes about 3.5 time longer to calculate the increment than the sind function. For increment 0.1 the sind function was faster by a factor of 3.5474860335195526. For increment 0.01 sind was faster by a factor of 3.3409598214285716. For 0.0001 the sind function was faster by a factor of 3.6095140273253272. For increment 1.0e-5 the sind function was faster by a factor of 3.392817738171016. For increment 1.0e-6 the sind function was faster by a factor of 3.2104690780000014. By looking at the trend of the elapsed times we can see that on average the function sind is much faster and as the increment get larger the gap between the two functions shrinks by a very small margin the smaller the increments become.